

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

Dr. Jie Zhang is currently an assistant professor of the school of Computer Science at Peking University, China. He is engaged in the research and design of storage systems, non-volatile memory and specialized processors. His research addresses the requirements for high-performance storage systems in the era of big data and artificial intelligence from the perspective of computer architecture. He is dedicated to breaking through the bottlenecks of data migration and the limitations of memory walls in the Von Neumann architecture. The projects he leads and have participated in have been funded by the Natural Science Foundation of China, Natural Science Foundation of Beijing, National Key Research and Development Program of China, Alibaba, Ant Research and Longsys Inc. with a cumulative total of over 4 million dollars. He has published over 80 papers in top-tier conferences and journals, including ISCA, OSDI, SOSP, HPCA, MICRO, ASPLOS, FAST, ATC, and Eurosys. **Among post-90s scholars in computer architecture and operating system, he ranks first globally in the number of papers published in top conferences (CSRankings, 2016–2026).** He was awarded NSFBJ Distinguished Young Scientist Fund (2026), ACM SIGCSE China Rising Star (2023), Intel Young Faculty Research Program (2023), NSFC Excellent Young Scientist Fund Overseas Program (2022), and OlympusMons Award (2025). He was inducted to HPCA Hall of Fame (2025). Please see more details from his personal website: <https://chaselab.wiki>.

WORK EXPERIENCE

Peking University, Beijing, China

School of Computer science

KAIST, Daejeon, Korea

Electrical Engineering (Computing Division)

Assistant Professor, PhD advisor

July 2021 –Current

Postdoctoral Researcher

March 2020 –May 2021

EDUCATION

Yonsei University, Incheon, Korea

PhD in Engineering

Advisor: Dr. Myoungsoo Jung

August 2015 – Feb 2020

University of Texas at Dallas, Richardson, Texas

PhD in Computer Engineering (transfer to Korea)

Advisor: Dr. Myoungsoo Jung

August 2014 – August 2015

University of Texas at Dallas, Richardson, Texas

Master of Science in Electrical Engineering

Advisor: Dr. Myoungsoo Jung

August 2012 – May 2014

Nanjing University of Posts and Telecommunications, Nanjing, China

BS in Communication Engineering (computer communication)

September 2008 – July 2012

PUBLICATIONS

2027

Eurosys

Beyond Overlap: Hierarchical Graph Abstraction for Out-of-Core Vector Search

CCF-A 会议

Lizheng Chen, Pinhuan Huang, Shaonan Ma, **Jie Zhang**, Qing Wang, Zhengyi Yang, Heng Zhang, Mingxing Zhang, Fan Yang, Guihai Chen, Chen Tian, Chengying Huan

22th ACM European Conference on Computer Systems

2026

TACO

Cambricon-FlexLLM: A Flexible Chiplet-Based Hybrid Architecture for On-Device 70B LLM Inference

CCF-A 期刊

Zhongkai Yu, Shengwen Liang, Tianyu Ma, Yunke Cai, Ziyuan Nan, Di Huang, Xinkai Song,

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

Yifan Hao, [Jie Zhang](#), Tian Zhi, Yongwei Zhao, Zidong Du, Xing Hu Qi Guo, Tianshi Chen

Transactions on Architecture and Code Optimizations

MICRO

LUTHE: Breaking the Memory Wall of High-Precision Encrypted LLM Nonlinearities via GPU–CSD Co-Design

CCF-A 会议

Xiurui Pan, Guoci Chen, Guangyu Sun, Mingzhe Zhang, [Jie Zhang*](#)

The 58th Annual IEEE/ACM International Symposium on Microarchitecture

SC

TPage: Transparent Page Management for CXL-Based Tiered Memory Systems

Top Conference

Maixin Ye, Suzhen Wu, [Jie Zhang](#), Bo Mao

International Conference for High Performance Computing, Networking, Storage and Analysis

ICS

Latency-SLO-Aware Memory Offloading for Large Language Model Inference

Top Conference

Chenxiang Ma, Hanyu Zhao, Zhisheng Ye, Zehua Yang, Tianhao Fu, Jiaxun Han, [Jie Zhang](#), Yingwei Luo, Xiaolin Wang, Zhenlin Wang, Yong Li, Diyu Zhou

ACM International Conference on Supercomputing

HPDC

BCCE: Block-Centric GPU Co-Design for Real-Time Range-Top-K Query at Scale

Top Conference

Chengying Huan, Ziheng Meng, Zhengyi Yang, Yongchao Liu, [Jie Zhang](#), Qing Wang, Jing Wang, Shaonan Ma, Zhibin Wang, Mingxing Zhang, Rong Gu, Baokun Wang, Guihai Chen, Chen Tian

ACM International Symposium on High-Performance Parallel and Distributed Computing

OSDI

Espresso: Constructing Cost-Efficient CXL JBOF via Inter-SSD Computing Resource

Top Conference

Sharing

Shushu Yi, Yuda An, Li Peng, Xiurui Pan, Qiao Li, Jieming Yin, Guangyan Zhang, Wenfei Wu, Chenxi Wang, Diyu Zhou, Zhenlin Wang, Xiaolin Wang, Yingwei Luo, Ke Zhou, [Jie Zhang*](#)

USENIX Symposium on Operating Systems Design and Implementation

OSDI

Merlin: An Efficient Adaptive Cache Eviction Algorithm via Fine-Grained Characterization

Top Conference

Liujia Li, Jinhao Guo, Yi Fan, Jianyu Wu, Zhenling Wang, [Jie Zhang](#), Yuval Tamir, Xiaolin Wang, Yingwei Luo, Diyu Zhou

USENIX Symposium on Operating Systems Design and Implementation

OSDI

MultiLane: Eliminating Centralized Bottlenecks in User-space Network Stack

Top Conference

Kang Hu, Shuqi Dong, Chuandong Li, Ran Yi, Zonghao Zhang, [Jie Zhang](#), Yingwei Luo, Xiaolin Wang, Zhenling Wang, Diyu Zhou

USENIX Symposium on Operating Systems Design and Implementation

ICLR

LouisKV: Efficient KV Cache Retrieval for Long Input-Output Sequences

Top Conference

Wenbo Wu, Qingyi Si, Xiurui Pan, Ye Wang, [Jie Zhang*](#)

The Fourteenth International Conference on Learning Representations

DAC

WILL: Write Invalidation Skipping for QLC NAND by Leveraging Data Lifespan and Latency Awareness

Top Conference

Beichen Ning, Yina Lv, Tianyu Ren, Bin Gao, Xiangyu Yao, [Jie Zhang](#), Qiao Li and Chun Jason

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

Xue

Design Automation Conference

DAC

Cachence: Fine-Grained Cache Partitioning in Both Time and Space

Top Conference

Liuja Li, Yuanlong Li, Yiming Yao, Jianyu Wu, Yi Fan, Jinhao Guo, Liren Zhu, [Jie Zhang](#), Xiaolin Wang, Yingwei Luo, Zhenlin Wang, Diyu Zhou

Design Automation Conference

FAST

Xerxes: Extensive Exploration of Scalable Hardware Systems with CXL-Based Simulation

Top Conference

Framework

Yuda An, Shushu Yi, Bo Mao, Qiao Li, Mingzhe Zhang, Diyu Zhou, Ke Zhou, Nong Xiao, Guangyu Sun, Yingwei Luo, [Jie Zhang*](#)

The USENIX Conference on File and Storage Technologies

HPCA

AutoGNN: End-to-End Hardware-Driven Graph Preprocessing for Enhanced GNN

Top Conference

Performance

Seungkwan Kang, Seungjun Lee, Donghyun Gouk, Miryeong Kwon, Hyunkyu Choi, Junhyeok Jang, Sangwon Lee, Huiwon Choi, [Jie Zhang](#), Wonil Choi, Mahmut T Kandemir, Myoungsoo Jung

IEEE International Symposium on High-Performance Computer Architecture

Eurosys

ColdCode: Cold Data Encoding for Enhanced Reliability and Lifetime in 3D NAND Flash

Top Conference

Qiao Li, Shangyu Wu, Zheng Wan, Yufei Cui, [Jie Zhang*](#), Chun Jason Xue

21th ACM European Conference on Computer Systems

2025

TCAD

CXL-DMSim: A Full-System CXL Disaggregated Memory Simulator With Comprehensive Silicon Validation

Top Journal

Yanjing Wang, Lizhou Wu, Wentao Hong, Yang Ou, Zicong Wang, Sunfeng Gao, [Jie Zhang](#), Sheng Ma, Dezun Dong, Xingyun Qi, Mingche Lai, Nong Xiao

PACT

Exploring Memory Tiering Systems in the CXL Era via FPGA-based Emulation and Device-

Top Conference

Side Management

Yiqi Chen, Xiping Dong, Zhe Zhou, Zhao Wang, [Jie Zhang](#), Guangyu Sun,

International Conference on Parallel Architecture and Compilation Techniques

NEWTON

Ultrafast and Reliable Domain-Wall and Skyrmion Logic in a Chirally Coupled Ferrimagnet

Top Journal

Yifei Ma, Dihua Wu, Fengbo Yan, Xiaoxiao Fang, Peixin Qin, Leran Wang, Li Liu, Laichuan Shen, Zhiqi Liu, Wenyun Yang, [Jie Zhang](#), Yanglong Hou, Yan Zhou, Feng Luo, Jinbo Yang, Kaiming Cai, Shuai Ning, Zhaochu Luo

Cell Press NEWTON

SOSP

Aeolia: Fast and Secure Userspace Interrupt-Based Storage Stack

Top Conference

Chuandong Li, Ran Yi, Zonghao Zhang, Jing Liu, Changwoo Min, [Jie Zhang](#), Yingwei Luo, Xiaolin Wang, Zhenling Wang, Diyu Zhou

ACM Symposium on Operating Systems Principles

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

HotStorage	SPDK+: Low Latency or High Power Efficiency? We Take Both <i>Endian Li, Shushu Yi, Li Peng, Qiao Li, Diyu Zhou, Zhenlin Wang, Xiaolin Wang, Bo Mao, Yingwei Luo, Ke Zhou, Jie Zhang*</i> <i>The 17th ACM Workshop on Hot Topics in Storage and File Systems</i>
ICDE	BL-Tree: The Best of Both Worlds by combining B+-Tree on Top and LSM-Tree on Bottom
Top Conference	<i>Maoxin Ye, Zuocheng Wang, Shengzhe Wang, Jiahong Chen, Chunfeng Du, Jie Zhang, Ke Zhou, Suzhen Wu, Bo Mao</i> <i>IEEE International Conference on Data Engineering</i>
ToS	Enhancing the Performance of Next-Generation SSD Arrays: A Holistic Approach
Top Journal	<i>Jie Zhang*, Shushu Yi, Xiurui Pan, Qiao Li, Qiang Li, Chenxi Wang, Bo Mao, Myoungsoo Jung</i> <i>ACM Transactions on Storage</i>
TCAD	FlexHMB: A Flexible HMB Design towards Bufferless Mobile Flash
Top Journal	<i>Yong Peng, Li Peng, Shaocong Sun, Lujia Yin, Qiao Li, Jie Zhang*</i> <i>ACM Transactions on Computer-Aided Design of Integrated Circuits and Systems</i>
ISCA	XHarvest: Rethinking High-Performance and Cost-Efficient SSD Architecture with
Top Conference	CXL-Driven Harvesting <i>Li Peng, Wenbo Wu, Shushu Yi, Xianzhang Chen, Chenxi Wang, Shengwen Liang, Zhe Wang, Nong Xiao, Qiao Li, Mingzhe Zhang, Jie Zhang*</i> <i>IEEE/ACM International Symposium on Computer Architecture</i>
ISCA	ArtMem: Adaptive Migration in Reinforcement Learning-Enabled Tiered Memory
Top Conference	<i>Xinyue Yi, Hongchao Du, Yu Wang, Jie Zhang, Qiao Li, Chun Jason Xue</i> <i>IEEE/ACM International Symposium on Computer Architecture</i>
Eurosys	Daredevil: Rescue Your Flash Storage from Inflexible Kernel Storage Stack
Top Conference	<i>Junzhe Li, Ran Shu, Jiayi Lin, Qingyu Zhang, Ziyue Yang, Jie Zhang, Yongqiang Xiong, Chenxiong Qian</i> <i>20th ACM European Conference on Computer Systems</i>
NSDI	Beehive: A Scalable Disaggregated Memory Runtime Exploiting Asynchrony of
Top Conference	Multithreaded Programs <i>Quanxi Li, Hong Huang, Ying Liu, Yanwen Xia, Jie Zhang, Mosong Zhou, Xiaobing Feng, Huimin Cui, Quan Chen, Yizhou Shan, Chenxi Wang</i> <i>22nd USENIX Symposium on Networked Systems Design and Implementation</i>
HPCA	InstAttention: In-Storage Attention Offloading for Cost-Effective Long-Context LLM
Top Conference	Inference <i>Xiurui Pan, Endian Li, Qiao Li, Shengwen Liang, Yizhou Shan, Ke Zhou, Yingwei Luo, Xiaolin Wang, Jie Zhang*</i> <i>IEEE International Symposium on High-Performance Computer Architecture</i>

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

HPCA	NeuVSA: A Unified and Efficient Accelerator for Neural Vector Search
Top Conference	Ziming Yuan, Lei Dai, Wen Li, Jie Zhang , Shengwen Liang, Ying Wang, Cheng Liu, Huawei Li, Jiafeng Guo, Peng Wang, Renhai Chen, Gong Zhang <i>IEEE International Symposium on High-Performance Computer Architecture</i>
HPCA	Criticality-Aware Instruction-Centric Bandwidth Partitioning for Data Center Applications
Top Conference	Liren Zhu, Liujia Li, Jianyu Wu, Yiming Yao, Zhan Shi, Jie Zhang , Zhenlin Wang, Xiaolin Wang, Yingwei Luo, Diyu Zhou <i>IEEE International Symposium on High-Performance Computer Architecture</i>
2024	
SOSP	BIZA: Design of Self-Governing Block-Interface ZNS AFA for Endurance and Performance
Top Conference	Shushu Yi, Shaocong Sun, Li Peng, Yingbo Sun, Ming-Chang Yang, Zhichao Cao, Qiao Li, Myoungsoo Jung, Ke Zhou, Jie Zhang* <i>ACM Symposium on Operating Systems Principles</i>
MICRO	NeoMem: Hardware/Software Co-Design for CXL-Native Memory Tiering
Top Conference	Zhe Zhou, Yiqi Chen, Tao Zhang, Yang Wang, Ran Shu, Shuotao Xu, Peng Cheng, Lei Qu, Yongqiang Xiong, Jie Zhang , Guangyu Sun <i>The 57th Annual IEEE/ACM International Symposium on Microarchitecture</i>
MICRO	Cambricon-Ilm: A chiplet-based hybrid architecture for on-device Inference of 70B LLM
Top Conference	Zhongkai Yu, Shengwen Liang, Tianyu Ma, Yunke Cai, Ziyuan Nan, Di Huang, Xinkai Song, Yifan Hao, Jie Zhang , Tian Zhi, Yongwei Zhao, Zidong Du, Xing Hu Qi Guo, Tianshi Chen <i>The 57th Annual IEEE/ACM International Symposium on Microarchitecture</i>
ToS	Extremely-Compressed SSDs with I/O Behavior Prediction
Top Journal	Xiangyu Yao, Qiao Li, Kaihuan Lin, Xinbiao Gan, Jie Zhang , Congming Gao, Zhirong Shen, Quanqing Xu, Chuanhui Yang, Jason Xue <i>ACM Transactions on Storage</i>
TC	Land of Oz: Resolving Orderless Writes in Zoned Namespace SSDs
Top Journal	Yingjia Wang, You Zhou, Fei Wu, Jie Zhang , Ming-Chang Yang <i>IEEE Transactions on Computers</i>
USENIX ATC	ScalaAFA: Constructing User-Space All-Flash Array Engine with Holistic Designs
Top Conference	Shushu Yi, Xiurui Pan, Qiao Li, Qiang Li, Chenxi Wang, Bo Mao, Myoungsoo Jung, Jie Zhang* <i>USENIX Annual Technical Conference</i>
USENIX ATC	ScalaCache: Scalable User-Space Page Cache Management with Software-Hardware
Top Conference	Coordination Li Peng, Yuda An, You Zhou, Chenxi Wang, Qiao Li, Chuanning Chen, Jie Zhang* <i>USENIX Annual Technical Conference</i>
Arxiv	From RDMA to RDCA: Toward High-Speed Last Mile of Data Center Networks Using Remote Direct Cache Access

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

Qiang Li, Qiao Xiang, Derui Liu, Yuxin Wang, Haonan Qiu, Xiaoliang Wang, [Jie Zhang](#), Ridi Wen, Haohao Song, Gexiao Tian, Chenyang Huang, Lulu Chen, Shaozong Liu, Yaohui Wu, Zhiwu Wu, Zicheng Luo, Yuchao Shao, Chao Han, Zongjie WU, Jianbo Dong, Zheng Cao, Jinbo WU, Jiwu Shu, Jiesheng Wu

Arxiv

TACO

Characterizing and Optimizing LDPC Performance on 3D NAND Flash Memories

Top Journal

Qiao Li, Yu Chen, Guanyu Wu, Yajuan Du, Min Ye, Xinbiao Gan, [Jie Zhang](#), Zhirong Shen, Jiwu Shu, Chun Jason Xue

ACM Transactions on Architecture and Code Optimization

ISCA

Flagger: Cooperative Acceleration for Large-Scale Cross-Silo Federated Learning

Top Conference

Aggregation

Xiurui Pan, Yuda An, Shengwen Liang, Bo Mao, Mingzhe Zhang, Qiao Li, Myoungsoo Jung, [Jie Zhang*](#)

IEEE/ACM International Symposium on Computer Architecture

ASPLOS

Achieving Near-Zero Read Retry for 3D NAND Flash Memory

Top Conference

Min Ye, Qiao Li, Yina Lv, [Jie Zhang](#), Tianyu Ren, Daniel Wen, Tei-Wei Kuo, Chun Jason Xue
ACM International Conference on Architectural Support for Programming Languages and Operating Systems

HPCA

StreamPIM: Streaming Matrix Computation in Racetrack Memory

Top Conference

Yuda An, Yunxiao Tang, Shushu Yi, Li Peng, Xiurui Pan, Guangyu Sun, Zhaochu Luo, Qiao Li, [Jie Zhang*](#)

IEEE International Symposium on High-Performance Computer Architecture

HPCA

BeaconGNN: Large-Scale GNN Acceleration with Asynchronous In-Storage Computing

Top Conference

Yuyue Wang, Xiurui Pan, Yuda An, [Jie Zhang*](#), Glenn Reinman*

IEEE International Symposium on High-Performance Computer Architecture

HPCA

LearnedFTL: A Learning-based Page-level FTL for Reducing Double Reads in Flash-based

Top Conference

SSDs

Shengzhe Wang, Zihang Lin, Suzhen Wu, Hong Jiang, [Jie Zhang](#), Bo Mao

IEEE International Symposium on High-Performance Computer Architecture

HPCA

Midas Touch: Invalid-Data Assisted Reliability and Performance Boost for 3D High-

Top Conference

Density Flash

QiaoLi, Hongyang Dang, Zheng Wan, Congming Gao, Min Ye, [Jie Zhang](#), Tei-Wei Kuo, Chun Jason Xue

IEEE International Symposium on High-Performance Computer Architecture

2023

NVMW

Optimizations of Linux Software RAID System for Next-Generation Storage

Shushu Yi, Yanning Yang, Yunxiao Tang, Zixuan Zhou, Junzhe Li, Yue Chen, Myoungsoo Jung,

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

*Jie Zhang**,

The 14th Annual Non-volatile Memories Workshop

SAC

BcBench: Exploring Throughput Processor Designs based on Blockchain Benchmarking

Xiurui Pan, Yue Chen, Shushu Yi, [Jie Zhang](#),*

The 38th ACM/SIGAPP Symposium on Applied Computing

CAL

Intelligent SSD Firmware for Zero-Overhead Journaling

SCI Journal

Hanyeoreum Bae, Donghyun Gouk, Seungjun Lee, Jiseon Kim, Sungjoon Koh, [Jie Zhang](#),
and Myoungsoo Jung*,*

IEEE Computer Architecture Letters (CAL)

2022

HotStorage

ScalaRAID: Optimizing Linux Software RAID System for Next-Generation Storage

*Shushu Yi, Yanning Yang, Yunxiao Tang, Zixuan Zhou, Junzhe Li, Yue Chen, Myoungsoo Jung,
[Jie Zhang*](#),*

14th USENIX Workshop on Hot Topics in Storage and File Systems (HotStorage 22)

THPC

Survey on Storage-Accelerator Data Movement

Zixuan Zhou, Shushu Yi, [Jie Zhang](#),*

CCF Transaction on High Performance Computing

NVMW

Integrating New Photonic-Based Heterogeneous Memory into Throughput Accelerators

[Jie Zhang](#), Myoungsoo Jung,

The 13th Annual Non-volatile Memories Workshop

NVMW

HAMS: Hardware Automated Memory-over-Storage for Large-scale Memory Expansion

[Jie Zhang](#), Miryeong Kwon, Donghyun Gouk, Sungjoon Koh, Nam Sung Kim,

Mahmut Taylan Kandemir, Myoungsoo Jung

The 13th Annual Non-volatile Memories Workshop

2021

MICRO

Ohm-GPU: Integrating New Optical Network and Heterogeneous Memory into GPU

Top Conference

Multi-Processors

[Jie Zhang](#), Myoungsoo Jung,

The 54th Annual IEEE/ACM International Symposium on Microarchitecture

ISCA

Revamping Storage Class Memory with Hardware Automated Memory-Over-Storage

Top Conference

Solution

[Jie Zhang](#), Miryeong Kwon, Donghyun Gouk, Sungjoon Koh, Nam Sung Kim,

Mahmut Kandemir, Myoungsoo Jung,

The IEEE/ACM International Symposium on Computer Architecture

NVMW

Architecting Throughput Processors with New Flash

[Jie Zhang](#), Myoungsoo Jung,

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

The 12th Annual Non-volatile Memories Workshop

NVMW

DRAM-less Accelerator for Energy Efficient Data Processing

Jie Zhang, Gyuyoung Park, David Donofrio, John Shalf, Myoungsoo Jung

The 12th Annual Non-volatile Memories Workshop

NVMW

Manycore-Based Scalable SSD Architecture Towards One and More Million IOPS

Jie Zhang, Miryeong Kwon, Michael Swift, Myoungsoo Jung,

The 12th Annual Non-volatile Memories Workshop

2020

ISCA

ZnG: Architecting GPU Multi-Processors with New Flash for Scalable Data Analysis

Top Conference

Jie Zhang, Myoungsoo Jung,

The IEEE/ACM International Symposium on Computer Architecture

FAST

Scalable Parallel Flash Firmware for Many-core Architectures

Top Conference

Jie Zhang, Miryeong Kwon, Michael Swift, Myoungsoo Jung,

The 18th USENIX Conference on File and Storage Technologies

HPCA

DRAM-less: Hardware Acceleration of Data Processing with New Memory

Top Conference

Jie Zhang, Gyuyoung Park, David Donofrio, John Shalf, Myoungsoo Jung

26th IEEE International Symposium on High-Performance Computer Architecture

ISPASS

Data Direct I/O Characterization for Future I/O System Exploration

Major Conference *Mohammad Alian, Yifan Yuan, Jie Zhang, Ren Wang, Myoungsoo Jung, Nam Sung Kim*

The IEEE International Symposium on Performance Analysis of Systems and Software

CAL

FastDrain: Removing Page Victimization Overheads in NVMe Storage Stack

SCI Journal

Jie Zhang, Miryeong Kwon, Sanghyun Han, Nam Sung Kim, Mahmut Kandemir and Myoungsoo Jung

IEEE Computer Architecture Letters (CAL)

2019

HPCA

FUSE: Fusing STT-MRAM into GPUs to Alleviate Off-Chip Memory Access Overheads

Top Conference

Jie Zhang, Myoungsoo Jung, Mahmut Kandemir,

25th IEEE International Symposium on High-Performance Computer Architecture

IISWC

Faster than Flash: An In-Depth Study of System Challenges for Emerging Ultra-Low

Major Conference **Latency SSDs**

Sungjoon Koh, Junkyeok Jang, Changrim Lee, Miryeong Kwon, Jie Zhang, Myoungsoo Jung,

The 2019 IEEE International Symposium on Workload Characterization

DAC

FlashGPU: Placing New Flash Next to GPU Cores

Top Conference

Jie Zhang, Miryeong Kwon, Hyojong Kim, Hyesoon Kim, Myoungsoo Jung,

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

The 56th Design Automation Conference (DAC), 2019

TPDS

Exploring Fault-Tolerant Erasure Codes for Scalable All-Flash Array Clusters

Top Journal

Sungjoon Koh, [Jie Zhang](#), Miryeong Kwon, Jungyeon Yoon, David Donofrio, Nam Sung Kim, Myoungsoo Jung,

IEEE Transactions on Parallel and Distributed Systems (TPDS)

NVMW

Addressing Fast-Detrapping for Reliable 3D NAND Flash Design

Mustafa Shihab, [Jie Zhang](#), Myoungsoo Jung, Mahmut Kandemir,

10th Annual Non-Volatile Memories Workshop -- Nominated as Memorable Paper Award

KCC

Maximizing GPU Cache Utilization with Adjustable Cache Line Management

[Jie Zhang](#), Myoungsoo Jung,

Korean Computer Congress (KCC), 2019 -- Excellent Paper Award

2018

OSDI

FlashShare: Punching Through Server Storage Stack from Kernel to Firmware for

Top Conference **Ultra-Low Latency SSDs**

[Jie Zhang](#), Miryeong Kwon, Donghyun Gouk, Changlim Lee, Mohammad Alian, Myoungjun

Chun, Mahmut Kandemir, Nam Sung Kim, Jihong Kim, Myoungsoo Jung,

13th USENIX Symposium on Operating Systems Design and Implementation

MICRO

Amber: Enabling Precise Full-System Simulation with Detailed Modeling of All SSD

Top Conference **Resources**

Donghyun Gouk, Miryeong Kwon, [Jie Zhang](#), Sungjoon Koh, Wonil Choi, Nam Sung Kim, Mahmut Kandemir, Myoungsoo Jung,

The 51st Annual IEEE/ACM International Symposium on Microarchitecture

TACO

ReveNAND: A Fast-Drift Aware Resilient 3D NAND Flash Design

Top Journal

Mustafa Shihab, [Jie Zhang](#), Myoungsoo Jung, Mahmut Kandemir,

ACM Transactions on Architecture and Code Optimization (TACO), 2018

Eurosys

FlashAbacus: A Self-governing Flash-based Accelerator for Low-power Systems

Top Conference [Jie Zhang](#), Myoungsoo Jung,

The European Conference on Computer Systems (EuroSys), 2018

IPDPS

CIAO: Cache Interference-Aware Throughput-Oriented Architecture and Scheduling for

Major Conference **GPUs**

[Jie Zhang](#), Shuwen Gao, Nam Sung Kim, Myoungsoo Jung,

32nd IEEE International Parallel & Distributed Processing Symposium (IPDPS), 2018

2017

CAL

SimpleSSD: Modeling Solid State Drive for Holistic System Simulation

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

- Major Journal** Myoungsoo Jung, **Jie Zhang**, Ahmed Abulila, Miryeong Kwon, Narges Shahidi, John Shalf, Nam Sung Kim and Mahmut Kandemir,
IEEE Computer Architecture Letters (CAL), 2017
- IISWC** **Understanding System Characteristics of Online Erasure Coding on Scalable, Distributed**
Major Conference **and Large-Scale SSD Array Systems**
Sungjoon Koh, **Jie Zhang**, Miryeong Kwon, Jungyeon Yoon, David Donofrio, Nam Sung Kim, Myoungsoo Jung,
IEEE International Symposium on Workload Characterization (IISWC), 2017
- IISWC** **TraceTracker: Hardware/Software Co-Evaluation for Large-Scale I/O Workload**
Major Conference **Reconstruction**
Miryeong Kwon, **Jie Zhang**, Gyuyoung Park, Wonil Choi, David Donofrio, John Shalf, Mahmut Kandemir, Myoungsoo Jung,
IEEE International Symposium on Workload Characterization (IISWC), 2017
- NPC** **An In-depth Performance Analysis of Many-Integrated Core for Communication Efficient Heterogeneous Computing**
Jie Zhang, Myoungsoo Jung,
IFIP International Conference on Network and Parallel Computing (NPC), 2017
- IJPP** **Enabling Realistic Logical Device Interface and Driver for NVM Express Enabled Full**
Major Journal **System Simulations**
Donghyun Gouk, **Jie Zhang**, Myoungsoo Jung,
IFIP International Conference on Network and Parallel Computing (NPC) and Invited for International Journal of Parallel Programming (IJPP), 2017

2016

- HPCA** **DUANG: Fast and Lightweight Page Migration in Asymmetric Memory Systems**
Top Conference Hao Wang, **Jie Zhang**, Gieseok Park, Sharmila Shridhar, Myoungsoo Jung, Nam Sung Kim,
IEEE Symposium on High Performance Computer Architecture (HPCA), 2016
- ASBD** **A Study for Block-level I/O Trace Reconstruction on All-Flash Arrays**
Miryeong Kwon, **Jie Zhang**, Gyuyoung Park, Myoungsoo Jung,
Workshop on Architectures and Systems for Big Data (ASBD@ISCA), 2016
- NVMSA** **An In-Depth Study of Next Generation Interface for Emerging Non-Volatile Memories**
Major Conference Wonil Choi, **Jie Zhang**, Shuwen Gao, Jaesoo Lee, Myoungsoo Jung, Mahmut Kandemir,
IEEE Non-Volatile Memory Systems and Applications Symposium (NVMSA), 2016
- INFLOW** **ROSS: A Design of Read-Oriented STT-MRAM Storage for Energy-Efficient Non-Uniform Cache Architecture**
Jie Zhang, Miryeong Kwon, Chanyoung Park, Myoungsoo Jung, Songkuk Kim,

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

INFLOW **USENIX Workshop on Interactions of NVM/Flash with Operating Systems and Workloads**
Couture: Tailoring STT-MRAM for Persistent Main Memory
Mustafa Shihab, [Jie Zhang](#), Shuwen Gao, Josep Sloan, Myoungsoo Jung,
USENIX Workshop on Interactions of NVM/Flash with Operating Systems and Workloads

2015

ASBD **CoDEN: A Hardware/Software CoDesign Emulation Platform for SSD-Accelerated Near Data Processing**
[Jie Zhang](#), Damian Szmulewicz, Erick Macias, Myoungsoo Jung,
The 5th Workshop on Architecture and System for Big Data (ASBD), 2015

PACT **NVMMU: Direct Solid State Disk Access for GPU-Accelerated Data Processing**
Top Conference *[Jie Zhang](#), David Donofrio, John Shalf, Myoungsoo Jung,*
The 24th International Conference on Parallel Architecture and Compilation Techniques

ICCD **OpenNVM: An Open-Sourced FPGA-based NVM Controller for Low Level Memory**
Major Conference **Characterization**
[Jie Zhang](#), Gieseok Park, David Donofrio, Mustafa Shihab, John Shalf and Myoungsoo Jung,
The 33rd International Conference on Computer Design (ICCD), 2015

PACT-SRC **Integrating 3D Resistive Memory Cache into GPGPU for Energy-Efficient Data Processing**
[Jie Zhang](#), David Donofrio, John Shalf and Myoungsoo Jung,
International Conference on parallel Architecture and Compilation Techniques (PACT) –
ACM SRC 2nd Runner Award, 2015

FAST-WiP **Shared Non-Volatile Memory Cache for Energy-Efficient High Throughput GPU Computing**
[Jie Zhang](#) and Myoungsoo Jung,
USENIX Conference on File and Storage Technologies Working in Progress (FAST WiP), 2015

2014

HotStorage **Power, Energy, and Thermal Considerations in SSD-Based I/O Acceleration**
[Jie Zhang](#), Myoungsoo Jung,
6th USENIX Workshop on Hot Topics in Storage and File Systems (HotStorage 14), 2014

PATENTS

- “GPU 原生的高性能全闪阵列软硬件设计”，易舒舒，张杰 (pending)
- “一种面向大模型推理的旁路山层转换层的计算型存储设备及数据处理方法”，张杰，潘修睿，周可 (pending)
- “一种基于主机端资源借用的成本效益型固态硬盘及主机”，张杰，彭力，武文博，易舒舒，陈咸彰，王晨曦，梁胜文，王喆，李乔，张明喆 (pending)
- “基于硬件一致性互联实现盘间资源共享的全闪阵列”，易舒舒，张杰，安昱达，彭力，李乔，尹捷明，吴文斐，周迪宇，王振林，汪小林，罗英伟，周可 (pending)

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

-
- “基于软硬协同设计的全闪阵列”，张杰，易舒舒，潘修睿，徐一茗，李乔，李强，王晨曦，毛波，Myoungsoo Jung (pending)
 - “Memory controlling device and computing device including the same”, Myoungsoo Jung, Donghyun Gouk, Miryeong Kwon, Sungjoon Koh, Jie Zhang, America (US20190171566A1)
 - “Flash-based accelerator and computing device including the same”, Myoungsoo Jung, Jie Zhang, America (US10824341B2, US20180321859, US20170285968)
 - “基于闪存的加速器和包含其的计算设备”, Myoungsoo Jung, Jie Zhang, China (CN107291424)
 - “基于闪存的加速器及包括该加速器的计算设备”, Myoungsoo Jung, Jie Zhang, China (CN109460369)
 - “Resistance switching memory-based accelerator”, Myoungsoo Jung, Gyuyoung PARK, Jie Zhang, America (US20180321880A1)
 - “PARALLEL PROCESSING UNIT, COMPUTING DEVICE INCLUDING THE SAME, AND THREAD SCHEDULING METHOD”, Jie Zhang, Myoungsoo Jung, America (WO2018021620)
 - “MEMORY CONTROL APPARATUS AND COMPUTING DEVICE INCLUDING SAME”, JUNG MYOUNGSOO, GOUK DONGHYUN, KWON MIRYEONG, KOH SUNGJOON, 정명수, JIE ZHANG, 국동현, 권미령, 고성준 장지에, Korea (KR1020180126267)
 - “COMPUTING DEVICE, METHOD OF PROCESSING INPUT/OUTPUT REQUEST, AND RECORDING MEDIUM”, Jie Zhang, Myoungsoo Jung, Donghyun Gouk, Miryeong Kwon, Sungjoon Koh, America (pending)
 - “FLASH-BASED COPROCESSOR”, Jie Zhang, Myoungsoo Jung, America (pending)
 - “FLASH STORAGE DEVICE AND METHOD OF SCHEDULING PAGE VICTIMIZATION”, Jie Zhang, Myoungsoo Jung, America (pending)
-

External Activities

Journal Guest Editor

- IEEE MICRO Special Issue

Journal Paper Reviewer

- IEEE Transactions on Computer
- ACM Transactions on Storage
- ACM Transactions on Architecture and Code Optimization
- ACM Transactions on Computer Systems
- IEEE Transactions on Parallel and Distributed Systems
- IEEE Computer Architecture Letters
- IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems

Conference Program Committee Member

- Eurosys'27
 - APSys'26
 - OSDI'26
 - MICRO'25'24
 - ISCA'26'25'24
 - HPCA'27'26'25'24
 - ATC'25'23
 - ISPASS'26'25
 - USENIX ATC'23
 - ChinaSys'24'22
 - SAC'22
-

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

Invited Talks and Presentations

- Invited talk, “Storage Architecture Design in AI Era”, Nanjing University of Posts and Telecommunications, 2026
- Invited talk, “Progress in Near-Data Processing for AI”, CCF SYS, 2025
- Invited talk, “Storage Architecture Design in AI Era”, CCF China Storage Conference, 2025
- Invited talk, “Memory System Design based on Coherence Interconnects”, HPC China, 2025
- Invited talk, “Memory System Design based on Coherence Interconnects”, Compunet, 2025
- Invited talk, “Flagger: Cooperative Acceleration for Large-Scale Cross-Silo Federated Learning Aggregation”, CNCC, 2024
- Invited talk, “From Emerging Applications to New Storage Technology: A Full-Stack Exploration”, Young PI forum, 2024
- Invited talk, “Progress in Near-Data Processing”, CCF Computility, 2024
- Invited talk, “Progress in Near-Data Processing”, ASSIST, 2024
- Invited talk, “Progress in Near-Data Processing”, CCF-TCIST, 2024
- Invited talk, “Progress in Near-Data Processing”, Renmin University, 2023
- Invited talk, “Storage System: From Firmware to Architecture”, Xiamen University, 2022
- Invited talk, “Revisiting the Designs of Storage Class Memory”, CCF CHIP, 2022
- Invited talk, “ZnG: Architecting GPU Multi-Processors with New Flash for Scalable Data Analysis”, Intel Computational Storage Lab, 2020
- Presentation, “ZnG: Architecting GPU Multi-Processors with New Flash for Scalable Data Analysis”, ISCA, online, 2020
- Presentation, “DRAM-less: Hardware Acceleration of Data Processing with New Memory”, HPCA, San Diego, CA, 2020
- Presentation, “Scalable Parallel Flash Firmware for Many-core Architectures”, FAST, Santa Clara, CA, 2020
- Presentation, “FUSE: Fusing STT-MRAM into GPUs to Alleviate Off-Chip Memory Access Overheads”, HPCA, Washington DC, 2019
- Presentation, “FlashGPU: Placing New Flash Next to GPU Cores”, DAC, Las Vegas, NV, 2019
- Presentation, “Maximizing GPU Cache Utilization with Adjustable Cache Line Management”, Jeju, South Korea, 2019
- Presentation, “FlashShare: Punching Through Server Storage Stack from Kernel to Firmware for Ultra-Low Latency SSDs”, Carlsbad, CA, 2018
- Presentation, “FlashAbacus: A Self-governing Flash-based Accelerator for Low-power Systems”, Porto, Portugal, 2018
- Presentation, “CIAO: Cache Interference-Aware Throughput-Oriented Architecture and Scheduling for GPUs”, IPDPS, Vancouver, Canada, 2018
- Presentation, “An In-depth Performance Analysis of Many-Integrated Core for Communication Efficient Heterogeneous Computing”, NPC, Anhui, China, 2017
- Presentation, “ROSS: A Design of Read-Oriented STT-MRAM Storage for Energy-Efficient Non-Uniform Cache Architecture”, Inflow, Savannah, GA, 2016
- Presentation, “Couture: Tailoring STT-MRAM for Persistent Main Memory”, Inflow, Savannah, GA, 2016
- Presentation, “CoDEN: A Hardware/Software CoDesign Emulation Platform for SSD-Accelerated Near Data Processing”, ASBD, Portland, OR, 2015
- Presentation, “NVMMU: Direct Solid State Disk Access for GPU-Accelerated Data Processing”, PACT, San Francisco, CA, 2015

JIE ZHANG

tel: +8619910031394|jiez@pku.edu.cn

-
- Presentation, “Integrating 3D Resistive Memory Cache into GPGPU for Energy-Efficient Data Processing”, PACT SRC, San Francisco, CA, 2015
 - Presentation, “OpenNVM: An Open-Sourced FPGA-based NVM Controller for Low Level Memory Characterization”, ICCD, New York city, NY, 2015
 - Presentation, “Shared Non-Volatile Memory Cache for Energy-Efficient High Throughput GPU Computing”, FAST WiP, Santa Clara, CA, 2015
 - Presentation, “Power, Energy, and Thermal Considerations in SSD-Based I/O Acceleration”, HotStorage, Philadelphia, PA, 2014

Teaching Experience

- Advanced Computer Architecture (Fall’26, Fall’25, Fall’24, Fall’23)
- Computer Architecture (Fall’26, Fall’25, Fall’24, Fall’23, Fall’22)
- Introduction of Computer System (Fall’26, Fall’25, Fall’24, Fall’23, Fall’22)
- IIT 3002 Operating Systems (Fall’15, Fall’16)
- IIT 6036 Computer Organization and Design (Fall’15, Fall’16)
- IIT 7024 Advanced System Architecture (Spring’16)

Honors

- 2026: NSFBJ Distinguished Young Scientist Program
- 2025: HPCA Hall of Fame
- 2025: OlympusMons Award
- 2023: Intel Young Faculty Research Program
- 2023: ACM SIGCSE China Rising Star Award
- 2022: NSFC Excellent Young Scientist Fund Overseas Program
- 2021: Our storage-class memory research is selected as KAIST breakthrough 50 years
- 2020-2021: Korean BK21+ Scholarship
- 2019: Annual Non-Volatile Memories Workshop (NVMW) -- Nominated as Memorable Paper Award
- 2019: Korea Computer Congress (KCC) -- Best Presentation Paper Award
- 2015: ACM Student Research Competition 2nd Runner Award